

Runtime Governance for Enterprise AI

Fail-closed execution. Audit-grade receipts. Measurable reliability.

Enterprise AI does not fail on capability. It fails on control. Execution outruns governance. The result: compounding risk, audit exposure, and operational blast radius with no deterministic containment.

The Four-Layer Enforcement Stack

1. AUTHORITY GATE

Governance evaluated before state mutation. Fail-closed halt if authority cannot be verified.

2. IMMUTABLE RECEIPTS

Every state change cryptographically attested. Append-only receipt ledger. Non-repudiation.

3. DRIFT GUARD

Behavioral constraint across time. Authority decay. No standing execution paths.

4. GATED SUBSTRATE

Physical isolation at the execution layer. Capability removed, not restricted.

Readiness Scan Deliverables

- ✓ Control-plane gap map
- ✓ Failure-mode heatmap
- ✓ Evidence & attestation checklist
- ✓ Prioritized 30/60/90 hardening roadmap

What You Walk Away With

A scored assessment of your AI control plane: where enforcement exists, where it does not, and what to harden first. Mapped to your regulatory exposure, operational blast radius, and deployment friction.

Engagement Tiers

AI Readiness Sprint	AI Hardening Engagement
\$7,500 / 2 weeks Readiness Scan + scored gap map + prioritized 30-day action plan. Designed for teams evaluating AI control-plane maturity.	\$25,000–\$40,000 / 6–10 weeks Full 4-layer enforcement audit + control-plane retrofit + hardening implementation. Designed for teams moving agentic AI into production.

SCHEDULE YOUR READINESS SCAN

orionintelligenceagency.com/book | dan@orionapexcapital.com

© Orion Apex Capital | United States | If governance is not enforced at runtime, it is not governance.